

A RethinkDB driver for Haskell

Étienne Laurin

2014-11-13

Setup

```
$ cabal update
$ cabal install rethinkdb -j
$ ghci

> import Database.RethinkDB.NoClash
> import qualified Database.RethinkDB as R
> :set -XOverloadedStrings
> default (Datum, ReQL, Integer, String)

> h <- connect "localhost" 28015 def
> run h dbList
["test","muni"]
```

Download These Slides

<http://atnnn.github.io/rethinkdb/haskell-driver-2014-11-13.pdf>

<http://goo.gl/UOX4iD>

Sample Data

```
> run h $ table "routes" ! 1  
{"display_name":"18-46th Avenue","id":"18"}
```

```
> run h $ table "runs" ! 1  
{"display_name":"Inbound to Fisherman's Wharf",  
"stops":["15926","14882",...],  
"direction_name":"Inbound",  
"id":"08BX_IB",  
"route_id":"8BX"}
```

```
> run h $ table "stops" ! 1  
{"display_name":"Merchant Rd & Golden Gate Br.",  
"location":Point<[-122.47587,37.8066699]>,  
"id":"114776"}
```

Rendered




Figure : All Routes

Compared to JavaScript

`default` and `do` are keywords in Haskell.

- ▶ `default` was renamed to `handle`
- ▶ `do` was renamed to `apply`

Almost all other operations have the same name as in the JavaScript driver:

<http://rethinkdb.com/api/javascript/>

Different Syntax

In JavaScript:

```
r.expr({foo: "bar"})('foo')
```

In Haskell:

```
["foo" := "bar"] ! "foo"
```

- ▶ ! to access fields
- ▶ := to build objects
- ▶ expr can usually be omitted

Function Composition

In JavaScript:

```
r.table("runs")  
  .map(r.row("stops").count())  
  .sum()
```

In Haskell:

```
R.sum  
  . R.map (\row -> count (row!"stops"))  
  . table $ "runs"
```

$(.) :: (b \rightarrow c) \rightarrow (a \rightarrow b) \rightarrow a \rightarrow c$

$R.map :: (\text{Expr } a, \text{Expr seq}) \Rightarrow (\text{ReQL} \rightarrow a) \rightarrow \text{seq} \rightarrow \text{ReQL}$

Function Composition

In JavaScript:

```
r.table("runs")  
  .map(r.row("stops").count())  
  .sum()
```

In Haskell:

```
table "runs"  
# R.map (\row -> (row!"stops") # count)  
# R.sum  
  
(#) :: (Expr b, Expr a) => a -> (a -> b) -> ReQL
```

Function Composition

In JavaScript:

```
> r.table("runs").map(r.row("stops").count()).sum()  
  .run(conn).then(console.log)  
6485
```

In Haskell:

```
> run h $ R.sum . R.map (count . (!"stops")) $ table "runs"  
6485
```

Types

The ReQL type represents an operation that can be performed on the server.

```
> table "routes" # get "N" # (!"display_name") :: ReQL
get(table("routes"), "N")["display_name"]
> run h it
"N-Judah"
```

Types

The Expr typeclass can build queries from other types.

```
class Expr e where  
  expr :: e -> ReQL
```

```
table "routes" :: Table  
get "N" :: Expr s => s -> ReQL  
table "routes" # get "N" :: ReQL
```

Types

The ReQL type has a Num, Fractional and Floating instance, allowing it to overload the built-in math operators:

```
count (table "routes") / 2 + 1 :: ReQL
```

But it has no Ord instance. There are separate comparison operators in the R namespace:

```
> run h $ count (table "stops") > 3000
```

Error: Could not deduce (Ord ReQL) arising from a use of **of** **>**

```
> run h $ count (table "stops") R.> 3000
```

```
true
```

Types

Query results are stored in a Datum:

data Datum

= Null

| Bool Bool

| String Text

| Number Double

| Array Array

| Object Object

| Time ZonedDateTime

| Point LonLat

| Line Line

| Polygon Polygon

| Binary ByteString

The ToDatum and FromDatum typeclasses allow converting to and from this type.

Running Queries

The run function can convert the result of a query to many different types including lists, cursors, tuples or maps.

```
> run h $ table "stops" ! "display_name" :: IO [String]
["19th Ave & Holloway Ave",
 "Merchant Rd & Golden Gate Br.", ...]
```

```
> c <- run h $
  table "stops" ! "display_name" :: IO (Cursor String)
> next c
Just "19th Ave & Holloway Ave"
```

```
> run h $ table "stops" ! "display_name"
  # sample 2 :: IO (String, String)
("Geary Blvd & Laguna St", "City College Bookstore")
```

```
> run h $ table "stops" ! 1 :: IO (Map String Datum)
```

Optional Arguments

In JavaScript:

```
r.circle([-122.411017, 37.773589], 500)  
r.circle([-122.411017, 37.773589], 500, {unit: 'm'})
```

In Haskell:

```
circle [-122.411017, 37.773589] 500  
ex circle [unit Meter] [-122.411017, 37.773589] 500
```

For example:

```
> let heavybit = [-122.411017, 37.773589]  
> run h $ table "stops" # R.filter (\stop ->  
    ex circle [unit Meter] heavybit 200  
    # includes (stop!"location"))  
    # R.map (!"display_name")  
["Harrison St & 9th St", "Folsom St & 9th St"]
```


Aggregations

```
> run h $ table "runs" # group (!"route_id") count  
[{"group":"1","reduction":4},  
 {"group":"10","reduction":2}, ...]
```

```
> run h $ table "runs"  
  # group (!"route_id") count  
  # group (!"reduction") count  
[{"group":1,"reduction":1},  
 {"group":2,"reduction":76},  
 {"group":3,"reduction":1},  
 {"group":4,"reduction":3}]
```

```
> run h $ table "runs"  
  # group (!"route_id") (avg . R.map count . (!"stops"))  
[{"group":"1","reduction":25.25},  
 {"group":"10","reduction":66}, ...]
```

Multiple Aggregations

RethinkDB itself cannot perform multiple aggregations on the same group. The Haskell driver circumvents this limitation, allowing queries like this:

```
> run h $ table "runs"
  # group (!"route_id")
  ((\x -> [avg x, R.sum x, R.max x])
   . R.map count . (!"stops"))
[{"group":"1","reduction":[25.25,101,49]},
 {"group":"10","reduction":[66,132,67]}, ...]
```

Multiple Aggregations

The Haskell driver rewrites multiple aggregations into a single aggregation. Pretty-printing the query lets us examine what it looks like:

```
> expr $ mapReduce  
  ((\x -> [avg x, R.sum x, R.max x])  
   . R.map count . (!"stops"))
```

```
(\b -> ((\g -> [div(g[0][0], g[0][1]), g[1], g[2]]))(  
  reduce(map(b, (\h -> [  
    (((\d -> count(d)))(h["stops"]), 1),  
    ((\e -> count(e)))(h["stops"]),  
    ((\f -> count(f)))(h["stops"]))],  
  (\i j -> [  
    [add(i[0][0], j[0][0]), add(i[0][1], j[0][1]),  
    add(i[1], j[1]),  
    branch(gt(i[2], j[2]), i[2], j[2])]))]))))
```

Importing Data

```
> let api path = str . ( ++ path) $  
  "http://proximobus.appspot.com/agencies/sf-muni"  
  
> run h $ tableCreate "routes"  
> run h $ table "routes" #  
  insert (http (api "/routes.json") def ! "items")  
{ "inserted": 81 }  
  
> run h $ table "routes" # ex update [nonAtomic] (\route ->  
  http (api "/routes/" + (route!"id") + ".json") def)  
{ "replaced": 81 }  
  
> run h $ tableCreate "stops"  
> run h $ table "routes" ! "id" # forEach (\id ->  
  table "runs" # insert (  
    http (api "/routes/" + id + "/runs.json") def ! "items"))  
{ "inserted": 168 }
```

Importing Data

```
> run h $ createTable "stops"
> run h $ table "runs" # forEach (\run ->
  flip apply [
    http (api "/routes/" + (run!"route_id") + "/runs/"
      + (run!"id") + "/stops.json") def ! "items"]
  $ \stops -> expr [
    table "runs" # get (run!"id") #
    update (const ["stops" := stops!"id"]),
    table "stops" # insert stops])
{"inserted":3691}

> run h $ table "stops" # ex update [nonAtomic] (\stop -> [
  "latitude" := remove,
  "longitude" := remove,
  "location" := point (stop!"longitude") (stop!"latitude")])
{"replaced":3691}
```

Visualising Data

```
import Geodetics.Grid
import Geodetics.Geodetic
import Geodetics.TransverseMercator

project :: LonLat -> (Double, Double)
project (LonLat lon lat) = let
  gd lat lon =
    Geodetic (lat*~degree) (lon*~degree) (0*~meter) WGS84
  pos = gd lat lon
  sf_sw = gd 37.614775 (-122.522278)
  offset = GridOffset (0*~meter) (0*~meter) (0*~meter)
  pt = toGrid (mkGridTM sf_sw offset _1) pos
  convert f = fromRational $ toRational (f pt /~ meter)
in (convert eastings, convert northings)
```

Visualising Data

```
import Diagrams.Prelude
import Diagrams.Backend.SVG.CmdLine
import Data.Colour.SRGB

main = do
  h <- fmap (use "muni") $ R.connect "localhost" 28015 def
  runs <- run h $ flip R.map (table "runs") $ \run -> [
    (\r -> expr [r!"bg_color", r!"fg_color"]) 'R.apply'
    [get (run!"route_id") (table "routes")],
    run!"stops" # R.map (\stop ->
      table "stops" # get stop # (!"location"))]

  let lines = flip map runs $ \((bg, fg), pts) ->
    let line = fromVertices (map (p2 . project) pts)
    in (line # lineColor (sRGB24read fg) # lw medium) 'atop'
      (line # lineColor (sRGB24read bg) # lw thick)
  mainWith (mconcat lines :: Diagram B R2)
```


Rendered



Figure : All Routes

Benchmarks

point-get



lower bound estimate upper bound

OLS regression	319 μ s	342 μ s	364 μ s
R ² goodness-of-fit	0.944	0.964	0.981
Mean execution time	323 μ s	343 μ s	361 μ s
Standard deviation	49.1 μ s	62.0 μ s	80.8 μ s




Figure : Criterion Benchmarks

Documentation

<https://hackage.haskell.org/package/rethinkdb/>

<pre>skip :: (Expr n, Expr seq) => n -> seq -> ReQL</pre> <p>Drop elements from the head of a sequence.</p> <pre>>>> run h \$ skip 2 [1, 2, 3, 4] [3,4]</pre>	<p>Synopsis</p>	<pre>data Index = PrimaryKey Index Key map :: (Expr a, Expr b) => (ReQL -> b) -> a -> ReQL withFields :: Expr seq => [ReQL] -> seq -> ReQL concatMap :: (Expr a, Expr b) => (ReQL -> b) -> a -> ReQL orderBy :: Expr s => [ReQL] -> s -> ReQL asc :: ReQL -> ReQL desc :: ReQL -> ReQL skip :: (Expr n, Expr seq) => n -> seq -> ReQL limit :: (Expr n, Expr seq) => n -> seq -> ReQL slice :: (Expr a, Expr b, Expr c) => a -> b -> c -> ReQL indexesOf :: (Expr fun, Expr seq) => fun -> seq -> ReQL isEmpty :: Expr seq => seq -> ReQL union :: (Expr a, Expr b) => a -> b -> ReQL sample :: (Expr n, Expr seq) => n -> seq -> ReQL group :: (Expr group, Expr reduction, Expr seq) => (ReQL -> group) reduce :: (Expr a, Expr s) => (ReQL -> ReQL -> a) -> s -> ReQL reduce0 :: (Expr base, Expr seq, Expr a) => (ReQL -> ReQL -> a) -> distinct :: Expr s => s -> ReQL contains :: (Expr x, Expr seq) => x -> seq -> ReQL mapReduce :: (Expr reduction, Expr seq) => (ReQL -> reduction) -> count :: Expr a => a -> ReQL</pre>
<pre>limit :: (Expr n, Expr seq) => n -> seq -> ReQL</pre> <p>Limit the size of a sequence.</p> <pre>>>> run h \$ limit 2 [1, 2, 3, 4] [1,2]</pre>		
<pre>slice :: (Expr a, Expr b, Expr c) => a -> b -> c -> ReQL</pre> <p>Cut out part of a sequence</p> <pre>>>> run h \$ slice 2 4 [1, 2, 3, 4, 5] [3,4]</pre>		
<pre>indexesOf :: (Expr fun, Expr seq) => fun -> seq -> ReQL</pre> <p>The position in the sequence of the elements that match the predicate</p> <pre>>>> run h \$ indexesOf (match "ba.") [str "foo", "bar", "ba"] [1,2]</pre>		

Figure : Haddock Documentation

Questions?

- ▶ IRC: Freenode #RethinkDB
- ▶ Source code and Issue tracker:
<https://github.com/AtnNn/haskell-rethinkdb/>
- ▶ Downloads and documentation:
<https://hackage.haskell.org/package/rethinkdb/>

Questions?

- ▶ IRC: Freenode #RethinkDB
- ▶ Source code and Issue tracker:
<https://github.com/AtnNn/haskell-rethinkdb/>
- ▶ Downloads and documentation:
<https://hackage.haskell.org/package/rethinkdb/>
- ▶ Why are there no Monads in this presentation?